

Assignment 1

Machine Learning (SE-807)



Submitted By:-

Shahzeb Awan

Reg no:-

362585

Submitted To:-

Dr. Qasim Umar Khan

Date: - 21, March, 2022

Q: Take a 5-featured Multi-Linear Regression Problem having at least 1000 training examples, plot the cost function vs. number of iterations, after that show the surface 3D plot of Cost function having variable θ_0 and θ_1 , at the end also show the contour plot of the cost function.

I have taken a problem of climate **Temperature** prediction using other features in relation to it like, **Humidity**, **Wind direction**, **Atmospheric pressure**, **Wind-speed** and **Visibility**.

First of all I downloaded the dataset from Kaggle having link as follows: -

<https://www.kaggle.com/datasets/zakriarehman/weather-data-for-linear-regression>

Then I selected my feature for prediction as Temperature and selected other features as inputs.

After arranging in a useable format in excel as follows:-

| | A | B | C | D | E | F | G | H | I | J | K |
|---|------------|---|------------------------|---|----------------------|---|------------------|---|-----------------|---|---------------|
| 1 | Humidity , | | Wind_Bearing_degrees , | | Pressure_millibars , | | Wind_Speed_kmh , | | Visibility_km , | | Temperature_c |
| 2 | 0.92 , | | 130 , | | 1021.6 , | | 11.27 , | | 8.05 , | | -0.55555556 |
| 3 | 0.73 , | | 330 , | | 1017 , | | 20.93 , | | 16.1 , | | 21.11111111 |
| 4 | 0.97 , | | 193 , | | 1013.99 , | | 5.9731 , | | 14.9086 , | | 16.6 |
| 5 | 0.82 , | | 300 , | | 1031.59 , | | 3.22 , | | 16.1 , | | 1.6 |
| 6 | 0.6 , | | 116 , | | 1020.88 , | | 10.8836 , | | 9.982 , | | 2.194444444 |
| 7 | 0.32 , | | 190 , | | 1015.33 , | | 21.4613 , | | 10.3523 , | | 27.53888889 |

Then I converted this excel data file to text form for usage in my MATLAB Program as shown on next page: -

| | | | | | | | | | | |
|-------------------------------------|------|--------|------|---------|---|---------|---|---------|---|--------------|
| Shahzeb_Awan_data_weather - Notepad | | | | | | | | | | |
| File | Edit | Format | View | Help | | | | | | |
| 0.92 | , | 130 | , | 1021.6 | , | 11.27 | , | 8.05 | , | -0.555555556 |
| 0.73 | , | 330 | , | 1017 | , | 20.93 | , | 16.1 | , | 21.11111111 |
| 0.97 | , | 193 | , | 1013.99 | , | 5.9731 | , | 14.9086 | , | 16.6 |
| 0.82 | , | 300 | , | 1031.59 | , | 3.22 | , | 16.1 | , | 1.6 |
| 0.6 | , | 116 | , | 1020.88 | , | 10.8836 | , | 9.982 | , | 2.194444444 |
| 0.32 | , | 190 | , | 1015.33 | , | 21.4613 | , | 10.3523 | , | 27.53888889 |
| 0.84 | , | 170 | , | 1009.04 | , | 7.9695 | , | 11.1251 | , | 19.97777778 |
| 0.86 | , | 30 | , | 1009.6 | , | 14.49 | , | 15.134 | , | 11.11111111 |
| 0.73 | , | 351 | , | 1018.39 | , | 14.007 | , | 15.8263 | , | 8.405555556 |
| 0.81 | , | 320 | , | 1003.89 | , | 6.44 | , | 7.8568 | , | 1.7 |
| 0.88 | , | 141 | , | 1021.28 | , | 14.007 | , | 6.0214 | , | -2.222222222 |
| 0.6 | , | 204 | , | 1019.52 | , | 1.4168 | , | 15.8263 | , | 21.9 |
| 0.87 | , | 1 | , | 1015.92 | , | 11.0285 | , | 14.9086 | , | 17.10555556 |
| 0.73 | , | 297 | , | 1013.06 | , | 4.0733 | , | 9.7566 | , | 17.77222222 |
| 0.39 | , | 35 | , | 1025.59 | , | 7.6636 | , | 9.982 | , | 24.95 |
| 0.92 | , | 310 | , | 1024.3 | , | 3.22 | , | 3.4615 | , | -2.711111111 |
| 0.78 | , | 180 | , | 1018.76 | , | 4.83 | , | 9.982 | , | 18.88888889 |
| 0.82 | , | 300 | , | 1000 | , | 20.2097 | , | 15.8263 | , | 4.227777778 |

Main code:-

```

%% Load Data
data = load('Shahzeb_Awan_data_weather.txt');
X = data(:, 1:5); % input features
y = data(:, 6); % output column is number 6
m = length(y); % length of dataset

%% ===== Part 1: Feature Normalization =====
% Scale features and set them to zero mean
fprintf('Normalizing Features ...\n');

[X mu sigma] = featureNormalize(X); % This function is explained in this
%report on later pages in detail

% Add intercept term to X
X = [ones(m, 1) X];

%% ===== Part 2: Gradient Descent =====

fprintf('Running gradient descent ...\n');

% Choosing some alpha value
alpha = 0.01;
num_iters = 1000;

% Init Theta and Running Gradient Descent

theta = [1;44;9;2;5;1]; % nice then zeros (comparitively better)
% gradient decent function is also explained later on
[theta, J1] = gradientDescentMulti2(X, y, theta, alpha, num_iters);

```

```

% Plot the convergence graph
figure;
plot(1:numel(J1), J1, 'b');
xlabel('Number of iterations');
ylabel('Cost J');

% Display gradient descent's result
fprintf('Theta computed from gradient descent: \n');
fprintf(' %f \n', theta);
fprintf('\n');
%if Humidity is 0.88 , Wind bearing degrees 141, pressure in mb is 1021.28,
wind speed in km/h is 14.007
%and visibility in km is 6.0214 then Temperature should be round about :-2.22
C
Temperature=[1, (0.88-mu(1))/sigma(1), (141-mu(2))/sigma(2), (1021.28-
mu(3))/sigma(3), (14.007-mu(4))/sigma(4), (6.0214-mu(5))/sigma(5)]*theta;

% =====

fprintf(['Predicted Temperature ' ...
        '(using gradient descent):\n %f\n'], Temperature);

%% ===== Part 3: Normal Equations =====
%% Analytical solution through ordinary least squares

fprintf('Solving with normal equations...\n');

%% Load Data
data = csvread('Shahzeb_Awan_data_weather.txt');
X = data(:, 1:5);
y = data(:, 6);
m = length(y);

% Add intercept term to X
X = [ones(m, 1) X];

% Calculate the parameters from the normal equation
theta = normalEqn(X, y); %this is also explained separately in following pages

% Display normal equation's result
fprintf('Theta computed from the normal equations: \n');
fprintf(' %f \n', theta);
fprintf('\n');

%if Humidity is 0.88 , Wind bearing degrees 141, pressure in mb is 1021.28,
wind speed in km/h is 14.007
%and visibility in km is 6.0214 then Temperature should be round about :-2.22
C

Temperature = [1,0.88 , 141 , 1021.28 , 14.007 , 6.0214]*theta;

% =====

```

```

fprintf(['Predicted Temperature ' ...
        '(using normal equations):\n %f\n'], Temperature);

%% ===== Part 4: Visualizing J(theta_0, theta_1) =====
fprintf('Visualizing J(theta_0, theta_1) ...\n')

% Grid over which we will calculate J
theta0_vals = linspace(-60, 100, 100);
theta1_vals = linspace(-60, 60, 100);

% initialize J_vals to a matrix of 0's
J_vals = zeros(length(theta0_vals), length(theta1_vals));

% Fill out J_vals
for i = 1:length(theta0_vals)
    for j = 1:length(theta1_vals)
        t = [theta0_vals(i); theta1_vals(j); 0.002519; -0.003176; -
0.179258; 0.361726];
        J_vals(i,j) = computeCostMulti(X, y, t);
    end
end

% Because of the way meshgrids work in the surf command, we need to
% transpose J_vals before calling surf, or else the axes will be flipped
J_vals = J_vals';
% Surface plot
figure;
surf(theta0_vals, theta1_vals, J_vals)
xlabel('\theta_0'); ylabel('\theta_1');

% Contour plot
figure;
% Plot J_vals as 15 contours spaced logarithmically between 0.01 and 100
contour(theta0_vals, theta1_vals, J_vals, logspace(-2, 3, 20))
xlabel('\theta_0'); ylabel('\theta_1');
hold on;
plot(theta(1), theta(2), 'rx', 'MarkerSize', 10, 'LineWidth', 2);

%my minimum cost Function Result at my Selected thetas from gradient decent:-
J=computeCostMulti(X, y, theta)

```

Gradient-decent function:-

```
function [theta, J_history] = gradientDescentMulti2(X, y, theta, alpha, num_iters)
% Initialize some useful values
m = length(y); % number of training examples
J_history = zeros(num_iters, 1);
for iter = 1:num_iters
% Perform a single gradient step on the parameter vector theta.
gradient=zeros(6,1);
    for i=1:m,
        for j=1:6,
            gradient(j,1)=gradient(j,1)+(theta'*X(i,:)'-y(i))*X(i,j);
        end
    end
theta=theta-alpha/m*gradient;
% Save the cost J in every iteration
J_history(iter) = computeCostMulti(X, y, theta);
end
end
```

Feature normalization:-

```
function [X_norm, mu, sigma] = featureNormalize(X)
% the mean value of each feature is 0 and the standard deviation
% is 1.
X_norm = X;
mu = zeros(1, 5);
sigma = zeros(1, 5);

mu=mean(X);
sigma=std(X);
X_norm=(X-mu)./sigma;
end
```

Closed Form Solution using normal Eqn:-

```
function [theta] = normalEqn(X, y)
%Computes the closed-form solution to linear regression
theta = zeros(5, 1);
theta=pinv(X'*X)*X'*y;
end
```

My Cost function:-

```
function J = computeCostMulti(X, y, theta)
m = length(y); % number of training examples
J = 0;
J=1/(2*m)*(X*theta-y)'+(X*theta-y);
end
```

Command Window Result:-

>>Shahzeb_Awan_Multi_linear_regression

Normalizing Features ...

Running gradient descent ...

Theta computed from gradient descent:

12.391067

-6.045475

0.255844

-0.367468

-1.279270

1.557708

Predicted Temperature (using gradient descent):

5.504655

Solving with normal equations...

Theta computed from the normal equations:

35.008717

-29.869259

0.002519

-0.003176

-0.179258

0.361726

Predicted Temperature (using normal equations):

5.502272

Visualizing J(theta_0, theta_1) ...

J =

23.1443




Figure 1 Cost function vs number of iterations

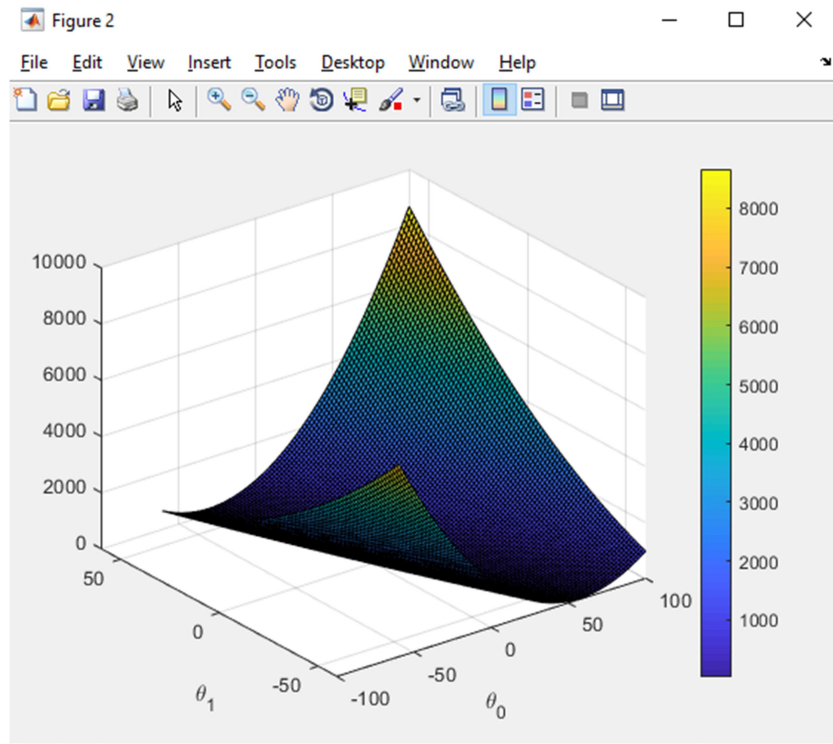


Figure 2 Cost Function vs variable Theta 0 and Theta 1

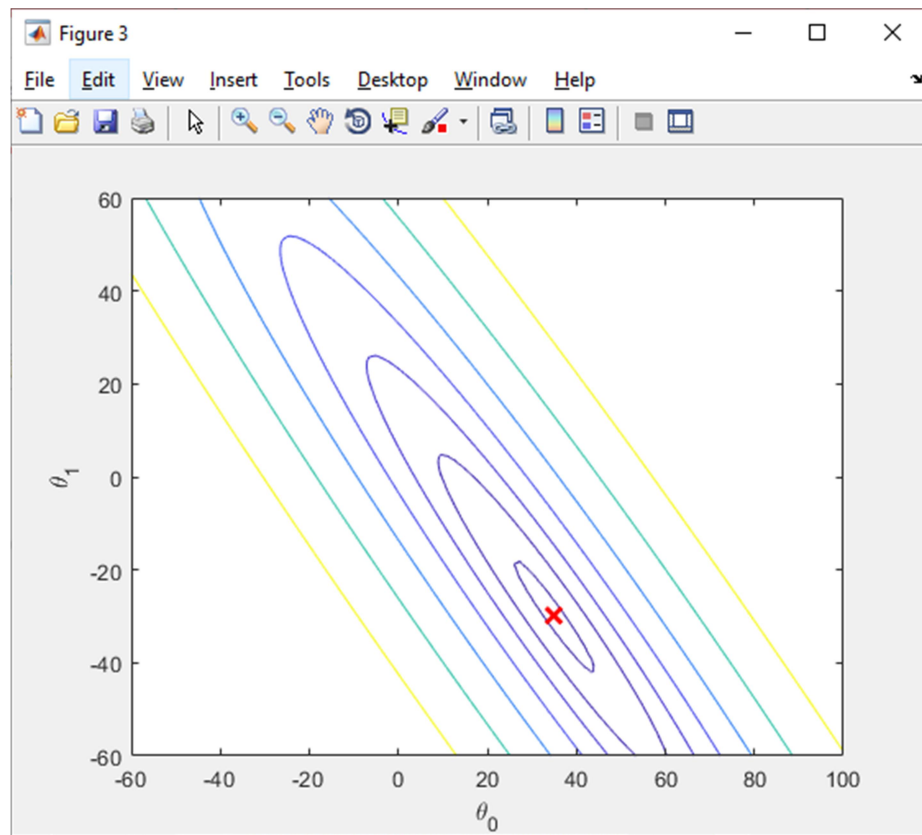


Figure 3 Contour plot of Cost function for theta 0 and theta 1

Following is my Github repository link for this project: -

<https://github.com/Shahzeb-Awan/Multi-Linear-Regression>

I have made this repository Private but after the marks allotment of this assignment I will make this repository public.

References:-

I have taken help from following two Github repositories and also through my Machine learning class knowledge: -

https://github.com/kk289/ML-Linear_Regression-MATLAB

<https://github.com/antoinevlt/Multivariate-linear-regression>